



Optimization algorithms

Quiz, 10 questions

✓ **Congratulations! You passed!**

Next Item



1 / 1
point

1.

Which notation would you use to denote the 3rd layer's activations when the input is the 7th example from the 8th minibatch?

☐ $a^{[3]\{7\}(8)}$

☐ $a^{[8]\{3\}(7)}$

☐ $a^{[8]\{7\}(3)}$

☒ $a^{[3]\{8\}(7)}$



Correct



1 / 1
point

2.

Which of these statements about mini-batch gradient descent do you agree with?

☐ Training one epoch (one pass through the training set) using mini-batch gradient descent is faster than training one epoch using batch gradient descent.

☒ One iteration of mini-batch gradient descent (computing on a single mini-batch) is faster than one iteration of batch gradient descent.



Correct

☐ You should implement mini-batch gradient descent without an explicit for-loop over different mini-batches, so that the algorithm processes all mini-batches at the same time (vectorization).



1 / 1
point

3.

Why is the best mini-batch size usually not 1 and not m, but instead something in-between?

☐ If the mini-batch size is 1, you end up having to process the entire training set before making any progress.



← **Un-selected is correct.**
Optimization algorithms
Quiz, 10 questions

☐ If the mini-batch size is 1, you lose the benefits of vectorization across examples in the mini-batch.

Correct

☐ If the mini-batch size is m , you end up with batch gradient descent, which has to process the whole training set before making progress.

Correct

☐ If the mini-batch size is m , you end up with stochastic gradient descent, which is usually slower than mini-batch gradient descent.


Un-selected is correct



1 / 1
point

4.

Suppose your learning algorithm's cost J , plotted as a function of the number of iterations, looks like this:



Which of the following do you agree with?

- ☐ Whether you're using batch gradient descent or mini-batch gradient descent, something is wrong.
- ☐ Whether you're using batch gradient descent or mini-batch gradient descent, this looks acceptable.
- ☒

Optimization algorithms

← Quiz, 10 questions

Correct

- ☐ If you're using mini-batch gradient descent, something is wrong. But if you're using batch gradient descent, this looks acceptable.

1 / 1
point

5.

Suppose the temperature in Casablanca over the first three days of January are the same:

Jan 1st: $\theta_1 = 10^\circ C$

Jan 2nd: $\theta_2 = 10^\circ C$

(We used Fahrenheit in lecture, so will use Celsius here in honor of the metric world.)

Say you use an exponentially weighted average with $\beta = 0.5$ to track the temperature: $v_0 = 0$, $v_t = \beta v_{t-1} + (1 - \beta)\theta_t$. If v_2 is the value computed after day 2 without bias correction, and $v_2^{\text{corrected}}$ is the value you compute with bias correction. What are these values? (You might be able to do this without a calculator, but you don't actually need one. Remember what is bias correction doing.)

- ☐ $v_2 = 10$, $v_2^{\text{corrected}} = 10$
- ☐ $v_2 = 10$, $v_2^{\text{corrected}} = 7.5$
- ☒ $v_2 = 7.5$, $v_2^{\text{corrected}} = 10$

Correct

- ☐ $v_2 = 7.5$, $v_2^{\text{corrected}} = 7.5$

1 / 1
point

6.

Which of these is NOT a good learning rate decay scheme? Here, t is the epoch number.

- ☐ $\alpha = \frac{1}{\sqrt{t}} \alpha_0$
- ☐ $\alpha = \frac{1}{1+2*t} \alpha_0$
- ☐ $\alpha = 0.95^t \alpha_0$
- ☒ $\alpha = e^t \alpha_0$

Correct




Optimization algorithms

point
Quiz, 10 questions

7.

You use an exponentially weighted average on the London temperature dataset. You use the following to track the temperature: $v_t = \beta v_{t-1} + (1 - \beta)\theta_t$. The red line below was computed using $\beta = 0.9$. What would happen to your red curve as you vary β ? (Check the two that apply)



☐ Decreasing β will shift the red line slightly to the right.



Un-selected is correct

☐ Increasing β will shift the red line slightly to the right.



Correct

True, remember that the red line corresponds to $\beta = 0.9$. In lecture we had a green line ($\beta = 0.98$) that is slightly shifted to the right.

☐ Decreasing β will create more oscillation within the red line.



Correct

True, remember that the red line corresponds to $\beta = 0.9$. In lecture we had a yellow line ($\beta = 0.98$) that had a lot of oscillations.

☐ Increasing β will create more oscillations within the red line.


Optimization algorithms

Un-selected is correct
Quiz, 10 questions



1 / 1
point

8.
Consider this figure:



These plots were generated with gradient descent; with gradient descent with momentum ($\beta = 0.5$) and gradient descent with momentum ($\beta = 0.9$). Which curve corresponds to which algorithm?

☒ (1) is gradient descent. (2) is gradient descent with momentum (small β). (3) is gradient descent with momentum (large β)



Correct

☐ (1) is gradient descent with momentum (small β), (2) is gradient descent with momentum (small β), (3) is gradient descent

☐ (1) is gradient descent. (2) is gradient descent with momentum (large β). (3) is gradient descent with momentum (small β)

☐ (1) is gradient descent with momentum (small β). (2) is gradient descent. (3) is gradient descent with momentum (large β)



1 / 1
point

9.
Suppose batch gradient descent in a deep network is taking excessively long to find a value of the parameters that achieves a small value for the cost function $\mathcal{J}(W^{[1]}, b^{[1]}, \dots, W^{[L]}, b^{[L]})$. Which of the following techniques could help find parameter values that attain a small value for \mathcal{J} ? (Check all that apply)



☒ Try using Adam



Correct

← Optimization algorithms



Try tuning the learning rate α

Correct



Try initializing all the weights to zero

Un-selected is correct



Try better random initialization for the weights

Correct



Try mini-batch gradient descent

Correct



1 / 1
point

10.

Which of the following statements about Adam is False?



Adam should be used with batch gradient computations, not with mini-batches.

Correct



We usually use "default" values for the hyperparameters β_1 , β_2 and ϵ in Adam ($\beta_1 = 0.9$, $\beta_2 = 0.999$, $\epsilon = 10^{-8}$)



Adam combines the advantages of RMSProp and momentum



The learning rate hyperparameter α in Adam usually needs to be tuned.

